

## This is the second part of the Telcom Churn Analysis

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: from sklearn.preprocessing import StandardScaler, RobustScaler, QuantileTransformer
from sklearn.feature_selection import SelectKBest, mutual_info_classif
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
```


```
In [3]: from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler
```

```
In [4]: df = pd.read_csv("/content/drive/MyDrive/DSC 550/data.csv")
df.head()
```

```
Out[4]:
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	Churn	gender_Male	Partner_Yes	Dependents_Yes	PhoneService	NonMult_PhoneService	...	NoStr
0	0	1	29.85	29.85	0	0	1	0	0		1	...
1	0	34	56.95	1889.50	0	1	0	0	1		0	...
2	0	2	53.85	108.15	1	1	0	0	1		0	...
3	0	45	42.30	1840.75	0	1	0	0	0		1	...
4	0	2	70.70	151.65	1	0	0	0	1		0	...

5 rows × 31 columns



```
In [5]: df['Churn'].value_counts()
```

```
Out[5]:
```

0	5163
1	1869

Name: Churn, dtype: int64

```
In [6]: X = df.drop('Churn', axis=1)
y = df['Churn']
```

```
In [7]: undersample = RandomUnderSampler(sampling_strategy='majority')
X_over, y_over = undersample.fit_resample(X, y)
```

```
In [8]: y_over.value_counts()
```

```
Out[8]:
```

0	1869
1	1869

Name: Churn, dtype: int64

```
In [9]: from sklearn.model_selection import train_test_split
```

```
In [10]: # Get X_train, X_test, y_train and y_test

X_train, X_test, y_train, y_test = train_test_split(X_over, y_over, test_size=0.20, random_state=42)
```

Setting up a pipeline

```
In [11]: pipe = Pipeline([('scaler', StandardScaler()),
                        ('selector', SelectKBest(mutual_info_classif, k=7)),
                        ('classifier', LogisticRegression())])
```

Define the search Space

```
In [12]:
```

```
search_space = [{'classifier':[LogisticRegression()],
                        'classifier__penalty':['l1','l2'],
                        'classifier__C':np.logspace(0,4,10)},
                {'classifier':[KNeighborsClassifier()],
                        'classifier__n_neighbors':[3,5,8,11],
                        'classifier__weights':['uniform','distance']},
                {'classifier':[RandomForestClassifier()],
                        'classifier__n_estimators':[10,100,1000],
                        'classifier__max_features':[1,2,3]}]
```

In [13]:

```
# Run the gridsearch
```

```
model = GridSearchCV(pipe,search_space,cv=10,verbose=0)
model.fit(X_train,y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_validation.py:372: FitFailedWarning:
100 fits failed out of a total of 370.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

-----

100 fits failed with the following error:

Traceback (most recent call last):

File "/usr/local/lib/python3.7/dist-packages/sklearn/model\_selection/\_validation.py", line 680, in \_fit\_and\_score

estimator.fit(X\_train, y\_train, \*\*fit\_params)

File "/usr/local/lib/python3.7/dist-packages/sklearn/pipeline.py", line 394, in fit

self.\_final\_estimator.fit(Xt, y, \*\*fit\_params\_last\_step)

File "/usr/local/lib/python3.7/dist-packages/sklearn/linear\_model/\_logistic.py", line 1461, in fit

solver = \_check\_solver(self.solver, self.penalty, self.dual)

File "/usr/local/lib/python3.7/dist-packages/sklearn/linear\_model/\_logistic.py", line 449, in \_check\_solver

% (solver, penalty)

ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

warnings.warn(some\_fits\_failed\_message, FitFailedWarning)

/usr/local/lib/python3.7/dist-packages/sklearn/model\_selection/\_search.py:972: UserWarning: One or more of the test scores are non-f

inite: [ nan 0.75183946 nan 0.74782609 nan 0.75150502

nan 0.75317726 nan 0.75150502 nan 0.74782609

nan 0.75351171 nan 0.75117057 nan 0.75217391

nan 0.75083612 0.7173913 0.70468227 0.73344482 0.71471572

0.74080268 0.71371237 0.75652174 0.71973244 0.70100334 0.72207358

0.71505017 0.70267559 0.71103679 0.71371237 0.69598662 0.71003344

0.71304348]

category=UserWarning,

Out[13]:

GridSearchCV(cv=10,

estimator=Pipeline(steps=[('scaler', StandardScaler()),

(('selector',

SelectKBest(k=7,

score\_func=<function mutual\_info\_classif at 0x7f14623fd0e0>)),

(('classifier', LogisticRegression()))],

param\_grid=[{'classifier': [LogisticRegression()],

'classifier\_\_C': array([1.00000000e+00, 2.78255940e+00, 7.74263683e+00, 2.15443469e+01,

5.99484250e+01, 1.668...6415883e+02, 1.29154967e+03,

3.59381366e+03, 1.00000000e+04]),

'classifier\_\_penalty': ['l1', 'l2']},

{'classifier': [KNeighborsClassifier(n\_neighbors=11)],

'classifier\_\_n\_neighbors': [3, 5, 8, 11],

'classifier\_\_weights': ['uniform', 'distance']},

{'classifier': [RandomForestClassifier()],

'classifier\_\_max\_features': [1, 2, 3],

'classifier\_\_n\_estimators': [10, 100, 1000]}])

Obtain predictions for train and test sets

In [14]:

```
best_estm = model.best_estimator_
print(best_estm)
```

Pipeline(steps=[('scaler', StandardScaler()),

(('selector',

SelectKBest(k=7,

score\_func=<function mutual\_info\_classif at 0x7f14623fd0e0>)),

(('classifier', KNeighborsClassifier(n\_neighbors=11))])

In [15]:

```
from sklearn.metrics import mean_squared_error
```

In [16]:

```
y_pred_train = best_estm.predict(X_train)
mse = mean_squared_error(y_pred_train,y_train)
print("MSE: %2f" %mse)
```

MSE: 0.219732

```
In [17]: y_pred_test = best_estm.predict(X_test)
mse = mean_squared_error(y_pred_test,y_test)
print("MSE: %2f" %mse)
```

MSE: 0.260695

Print classification Report

```
In [18]: from sklearn.metrics import confusion_matrix,classification_report
```

```
In [19]: cm = confusion_matrix(y_test,y_pred_test)
```

```
In [20]: print(cm)
```

```
[[269 110]
 [ 85 284]]
```

```
In [21]: print(classification_report(y_test,y_pred_test))
```

	precision	recall	f1-score	support
0	0.76	0.71	0.73	379
1	0.72	0.77	0.74	369
accuracy			0.74	748
macro avg	0.74	0.74	0.74	748
weighted avg	0.74	0.74	0.74	748

Having adjusted my class size using random undersampler, the classification of those that churn has improved from 55% to 75%. I also incorporated feature selection into the pipeline model to reduce unnecessary noise.

```
In [21]:
```